

```
In [1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.svm import SVR
from sklearn.feature_selection import SelectKBest, f_regression, RFE
```

```
In [2]: Data = pd.read_csv('CarPrice_Assignment.csv')
```


```
In [3]: df = pd.DataFrame(Data)
```

```
In [4]: df.head(2)
```

```
Out[4]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd


2 rows × 26 columns



```
In [5]: df.describe()
```

```
Out[5]:
```

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000
mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000



In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null    int64
1   symboling              205 non-null    int64
2   CarName               205 non-null    object
3   fueltype              205 non-null    object
4   aspiration             205 non-null    object
5   doornumber            205 non-null    object
6   carbody               205 non-null    object
7   drivewheel            205 non-null    object
8   enginelocation        205 non-null    object
9   wheelbase             205 non-null    float64
10  carlength             205 non-null    float64
11  carwidth              205 non-null    float64
12  carheight             205 non-null    float64
13  curbweight            205 non-null    int64
14  enginetype            205 non-null    object
15  cylindernumber        205 non-null    object
16  enginesize            205 non-null    int64
17  fuelsystem            205 non-null    object
18  boreratio             205 non-null    float64
19  stroke                205 non-null    float64
20  compressionratio      205 non-null    float64
21  horsepower            205 non-null    int64
22  peakrpm               205 non-null    int64
23  citympg               205 non-null    int64
24  highwaympg            205 non-null    int64
25  price                 205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
```

In [7]: `df.isnull().sum()`

```
Out[7]: car_ID      0
        symboling   0
        CarName     0
        fueltype    0
        aspiration   0
        doornumber   0
        carbody      0
        drivewheel   0
        enginelocation 0
        wheelbase    0
        carlength    0
        carwidth     0
        carheight    0
        curbweight   0
        enginetype   0
        cylindernumber 0
        enginesize    0
        fuelsystem   0
        boreratio    0
        stroke       0
        compressionratio 0
        horsepower   0
        peakrpm      0
        citympg      0
        highwaympg   0
        price        0
        dtype: int64
```

```
In [8]: df.duplicated().sum()
```

```
Out[8]: 0
```

```
In [9]: import warnings
        warnings.filterwarnings('ignore')
```

```
In [10]: num_df = df.select_dtypes(include=['int64', 'float64'])

        skewness = num_df.skew()
        kurtosis = num_df.kurt()
```


```
In [11]: skewness
```

```
Out[11]: car_ID          0.000000
        symboling      0.211072
        wheelbase      1.050214
        carlength       0.155954
        carwidth        0.904003
        carheight       0.063123
        curbweight      0.681398
        enginesize       1.947655
        boreratio       0.020156
        stroke         -0.689705
        compressionratio 2.610862
        horsepower      1.405310
        peakrpm         0.075159
        citympg         0.663704
        highwaympg      0.539997
        price           1.777678
        dtype: float64
```

```
In [12]: kurtosis
```


```
Out[12]: car_ID          -1.200000
        symboling      -0.676271
        wheelbase      1.017039
        carlength      -0.082895
        carwidth        0.702764
        carheight      -0.443812
        curbweight     -0.042854
        enginesize       5.305682
        boreratio      -0.785042
        stroke          2.174396
        compressionratio 5.233054
        horsepower      2.684006
        peakrpm         0.086756
        citympg         0.578648
        highwaympg      0.440070
        price           3.051648
        dtype: float64
```

```
In [13]: plt.figure(figsize=(8, 6))
        sns.histplot(df['price'], kde=True, bins=20, color='blue')
        plt.title('Distribution of Price (Histogram)', fontsize=14, color='red')
        plt.axvline(df['price'].mean(), color='red', linestyle='--', label='Mean')
        plt.axvline(df['price'].median(), color='green', linestyle='--', label='Median')
        plt.xlabel('Price', fontsize=12)
        plt.ylabel('Frequency', fontsize=12)
        plt.legend(fontsize=10)
        plt.show()
```

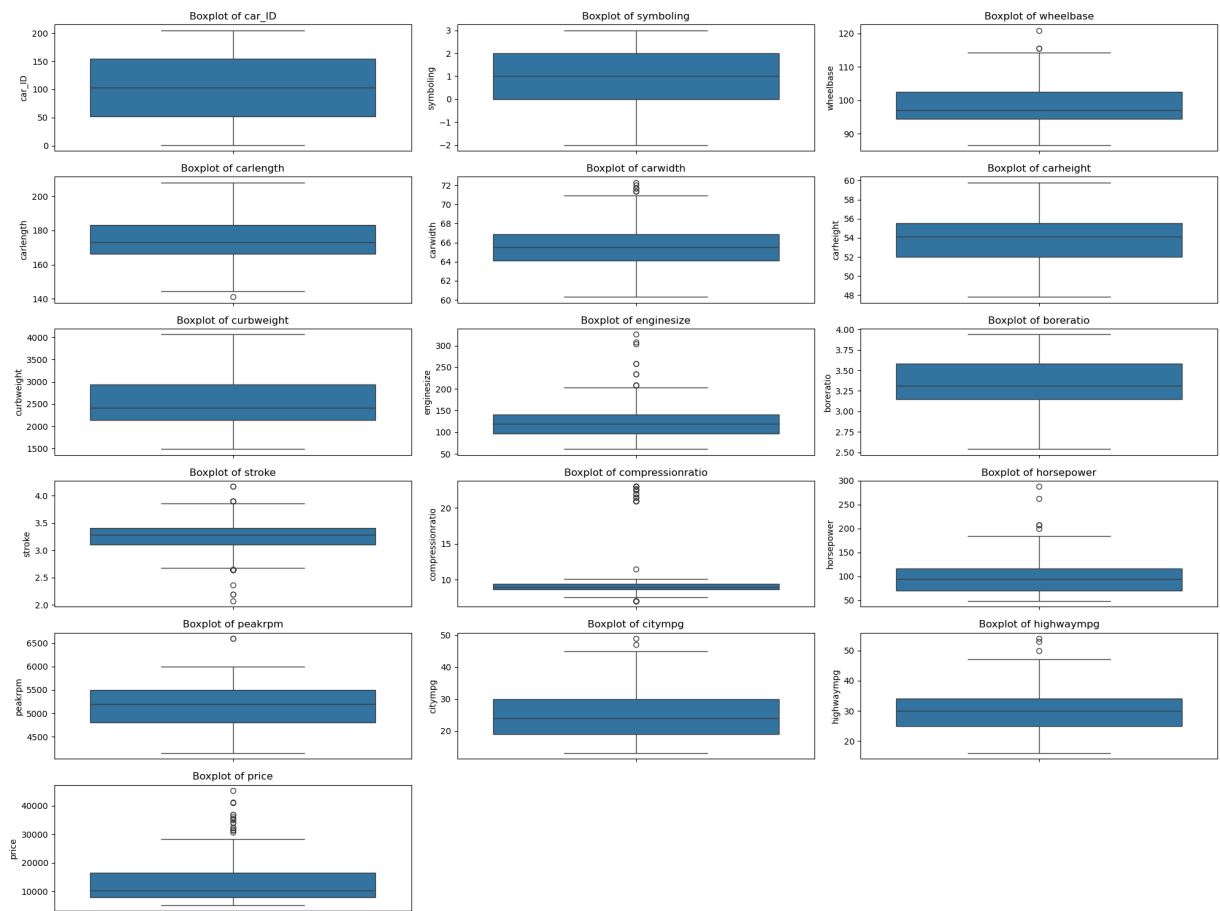


```
In [26]: correlation_matrix = num_df.corr()

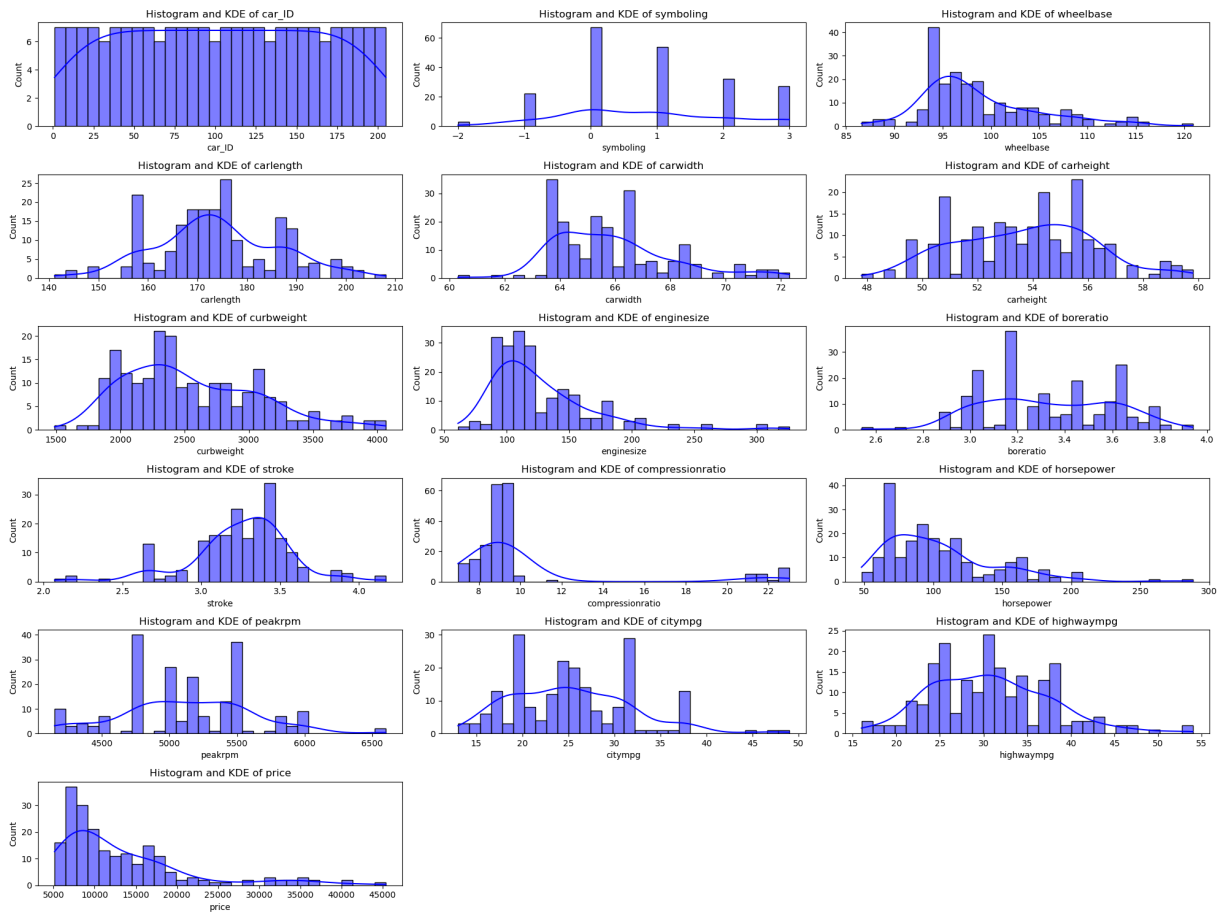
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Correlation Heatmap")
plt.show()
```



```
In [28]: # Select only numerical columns
numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns
# Set up the plot size
plt.figure(figsize=(20, 15))
# Loop through each numerical column and draw a boxplot
for i, column in enumerate(numerical_columns, 1):
    plt.subplot(len(numerical_columns) // 3 + 1, 3, i)
    sns.boxplot(data=df, y=column)
    plt.title(f"Boxplot of {column}")
    plt.tight_layout()
plt.show()
```



```
In [29]: # Select only numerical columns
numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns
# Set up the plot size
plt.figure(figsize=(20, 15))
# Loop through each numerical column and draw a histogram with KDE
for i, column in enumerate(numerical_columns, 1):
    plt.subplot(len(numerical_columns) // 3 + 1, 3, i)
    sns.histplot(data=df, x=column, kde=True, color='blue', bins=30)
    plt.title(f"Histogram and KDE of {column}")
    plt.tight_layout()
plt.show()
```




```
In [31]: df.head(10)
```



Out[31]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewhe
0	1	3	alfa-romero giulia	gas	std	two	convertible	rv
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rv
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rv
3	4	2	audi 100 ls	gas	std	four	sedan	fv
4	5	2	audi 100ls	gas	std	four	sedan	4v
5	6	2	audi fox	gas	std	two	sedan	fv
6	7	1	audi 100ls	gas	std	four	sedan	fv
7	8	1	audi 5000	gas	std	four	wagon	fv
8	9	1	audi 4000	gas	turbo	four	sedan	fv
9	10	0	audi 5000s (diesel)	gas	turbo	two	hatchback	4v

10 rows × 26 columns



In [32]:

```

# Identify numerical columns
numerical_columns = df.select_dtypes(include=['int64', 'float64']).columns
# Detect outliers using IQR for each column
for col in numerical_columns:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

# Count outliers
outliers = df[(df[col] < lower_bound) | (df[col] > upper_bound)]
print(f"Column: {col}, Outliers: {len(outliers)}")

```

```

Column: car_ID, Outliers: 0
Column: symboling, Outliers: 0
Column: wheelbase, Outliers: 3
Column: carlength, Outliers: 1
Column: carwidth, Outliers: 8
Column: carheight, Outliers: 0
Column: curbweight, Outliers: 0
Column: enginesize, Outliers: 10
Column: boreratio, Outliers: 0
Column: stroke, Outliers: 20
Column: compressionratio, Outliers: 28
Column: horsepower, Outliers: 6
Column: peakrpm, Outliers: 2
Column: citympg, Outliers: 2
Column: highwaympg, Outliers: 3
Column: price, Outliers: 15

```

```

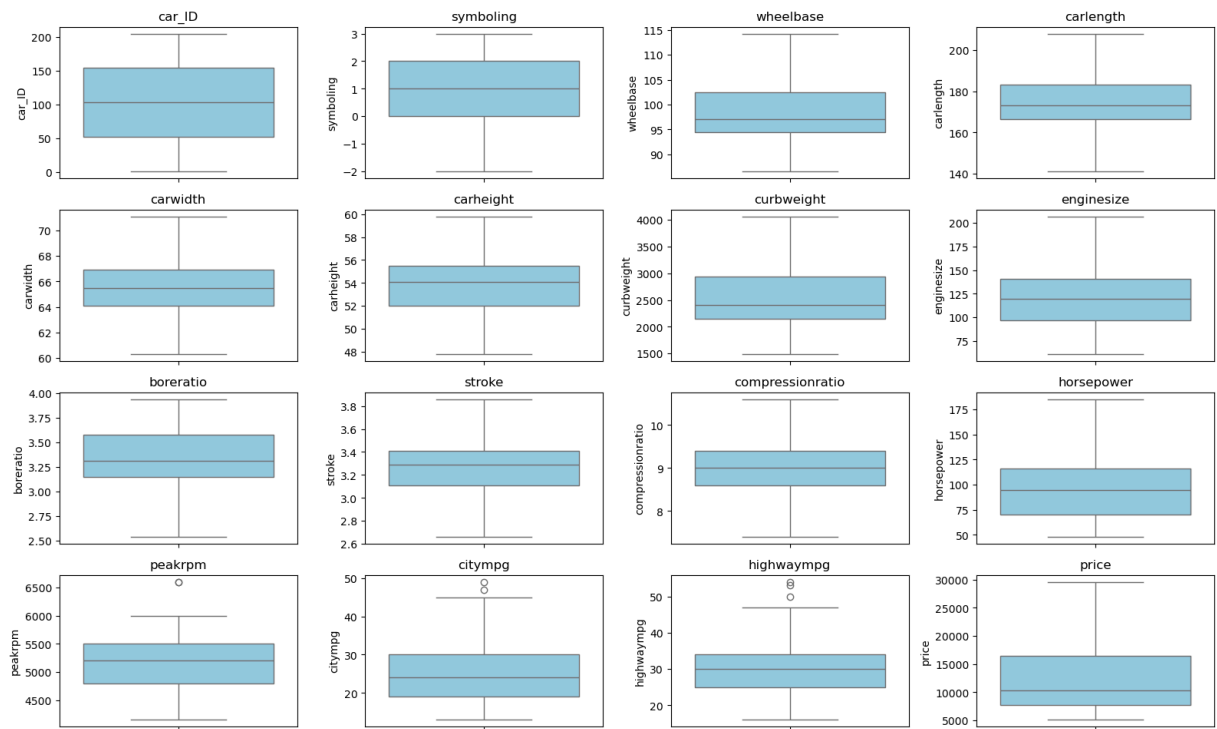
In [33]: # Create a cleaned dataset
df_cleaned = df.copy()
# List of columns to check for outliers
columns_with_outliers = ['wheelbase', 'carlength', 'carwidth', 'enginesize', 'stroke']
# Capping outliers
for column in columns_with_outliers:
    Q1 = df_cleaned[column].quantile(0.25)
    Q3 = df_cleaned[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    # Cap the outliers
    df_cleaned[column] = df_cleaned[column].apply(lambda x: lower_bound if x < lower_bound else

```

```

In [36]: # Set up the figure size
plt.figure(figsize=(16, 12))
# Create boxplots for all numeric columns in df_cleaned
for i, column in enumerate(df_cleaned.select_dtypes(include='number').columns, 1):
    plt.subplot(5, 4, i) # Create a grid of subplots (adjust rows/cols as needed)
    sns.boxplot(y=df_cleaned[column], color='skyblue')
    plt.title(column)
# Adjust layout to avoid overlap
plt.tight_layout()
plt.show()

```



In [ ]: `# Encoding`

```
In [37]: # Creating a new dataframe df_cleaned based on the original df
df_cleaned = df.copy()
# Label encoding for the 'symboling' column (ordinal data)
label_encoder = LabelEncoder()
df_cleaned['symboling'] = label_encoder.fit_transform(df_cleaned['symboling'])
# One-Hot Encoding for categorical columns (nominal data)
df_cleaned = pd.get_dummies(df_cleaned, columns=['fueltype', 'aspiration', 'doornum',
'carbody', 'drivewheel', 'enginelocation', 'wheelbase', 'carlength', 'carwidth', 'carh',
'enginetype', 'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke', 'co

# Check the first few rows to confirm the encoding
df_cleaned.head()
```

```
Out[37]:
```

	car_ID	symboling	CarName	curbweight	price	fueltype_diesel	fueltype_gas	aspi
0	1	5	alfa-romero giulia	2548	13495.0	False	True	
1	2	5	alfa-romero stelvio	2548	16500.0	False	True	
2	3	3	alfa-romero Quadrifoglio	2823	16500.0	False	True	
3	4	4	audi 100 ls	2337	13950.0	False	True	
4	5	4	audi 100ls	2824	17450.0	False	True	

5 rows × 556 columns

```
In [76]: numeric_columns = df_cleaned.select_dtypes(include=['int64', 'float64']).columns.to
numeric_columns.remove('price')
scaler = StandardScaler()
df_cleaned[numeric_columns] = scaler.fit_transform(df_cleaned[numeric_columns])
X = df_cleaned.drop(columns = ['price'])
y = df_cleaned['price']
```

```
In [78]: # Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta
```

```
In [80]: # Separate features (X) and target variable (y)
X = df_cleaned.drop(columns=['price']) # Features (all columns except 'price')
y = df_cleaned['price'] # Target variable ('price')
```

```
In [82]: selector = SelectKBest(score_func=f_regression,k=20)
x_new = selector.fit_transform(X,y)
```

```
In [84]: Selected_features = X.columns[selector.get_support()]
feature_score = pd.DataFrame({'features':X.columns,'score':selector.scores_}).sort_
```

```
In [86]: print("Selected Features:", Selected_features)
print("\nFeature Scores:")
print(feature_score)
```

```
Selected Features: Index(['curbweight', 'drivewheel_fwd', 'drivewheel_rwd', 'enginetype_ohc',
                        'enginetype_ohcv', 'cylindernumber_eight', 'cylindernumber_four',
                        'cylindernumber_six', 'enginesize_209', 'fuelsystem_2bbl',
                        'fuelsystem_mpf', 'boreratio_3.8', 'compressionratio_8.0',
                        'horsepower_182', 'horsepower_184', 'peakrpm_4750', 'citympg_14',
                        'citympg_16', 'highwaympg_16', 'highwaympg_25'],
                        dtype='object')
```

Feature Scores:

	features	score
2	curbweight	468.594431
249	cylindernumber_four	192.612277
16	drivewheel_rwd	140.059236
15	drivewheel_fwd	115.353549
303	fuelsystem_mpf	74.082624
..	...	...
382	compressionratio_7.5	0.000788
127	carlength_181.7	0.000774
443	horsepower_111	0.000100
227	carheight_55.7	0.000036
222	carheight_55.1	0.000006

[554 rows x 2 columns]

```
In [88]: X_selected = X[Selected_features]
X_train_selected, X_test_selected, y_train, y_test = train_test_split(X_selected, y
```

```
In [90]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_selected)
```

```
X_test_scaled = scaler.fit_transform(X_test_selected)
```

```
In [92]: # Initialize models
models={"Linear Regression":LinearRegression(),
        "Decision Tree Regressor":DecisionTreeRegressor(),
        "Random Forest Regressor":RandomForestRegressor(),
        "Gradient Boosting Regressor":GradientBoostingRegressor(),
        "Support Vector Regressor":SVR()}
```

```
In [94]: print("Training set shape (features):", X_train_scaled.shape)
print("Test set shape (features):", X_test_scaled.shape)
print("Training set shape (target):", y_train.shape)
print("Test set shape (target):", y_test.shape)
```

```
Training set shape (features): (164, 20)
Test set shape (features): (41, 20)
Training set shape (target): (164,)
Test set shape (target): (41,)
```

```
In [98]: # MODEL EVALUATION
results={} # use to store evaluation result

for model_name, model in models.items():
    # fit the model
    model.fit(X_train_scaled,y_train)
    # make the prediction
    y_pred = model.predict(X_test_scaled)
    # Evaluate the model
    mse = mean_squared_error(y_test,y_pred)
    mae = mean_absolute_error(y_test,y_pred)
    r2 = r2_score(y_test,y_pred)
    rmse = np.sqrt(mean_squared_error(y_test,y_pred))
    # Store the results
    results[model_name] = {"MSE": mse, "MAE":mae,"RMSE":rmse,"R²": r2,}
```

```
In [100... # Convert results to DataFrame for better visualization
results_df = pd.DataFrame(results).T
print(results_df)
```

	MSE	MAE	RMSE	R <sup>2</sup>
Linear Regression	9.542248e+06	2009.757898	3089.052863	0.879126
Decision Tree Regressor	1.518954e+07	2558.709341	3897.375512	0.807591
Random Forest Regressor	1.705659e+07	2498.787541	4129.962554	0.783941
Gradient Boosting Regressor	2.123520e+07	2517.967415	4608.166600	0.731009
Support Vector Regressor	8.670856e+07	5695.907738	9311.743165	-0.098355

```
In [102... # Finding Best model
```

```
In [104... best_model_name = max(results, key=lambda x: results[x]['R²'])
best_model = models[best_model_name]
print(f"The best model is: {best_model_name}")
```

```
The best model is: Linear Regression
```

```
In [106... # Feature selection
```

```
In [112... # Extract coefficients as feature importances
feature_importances = pd.Series(best_model.coef_, index=X.columns).sort_values(ascending=True)

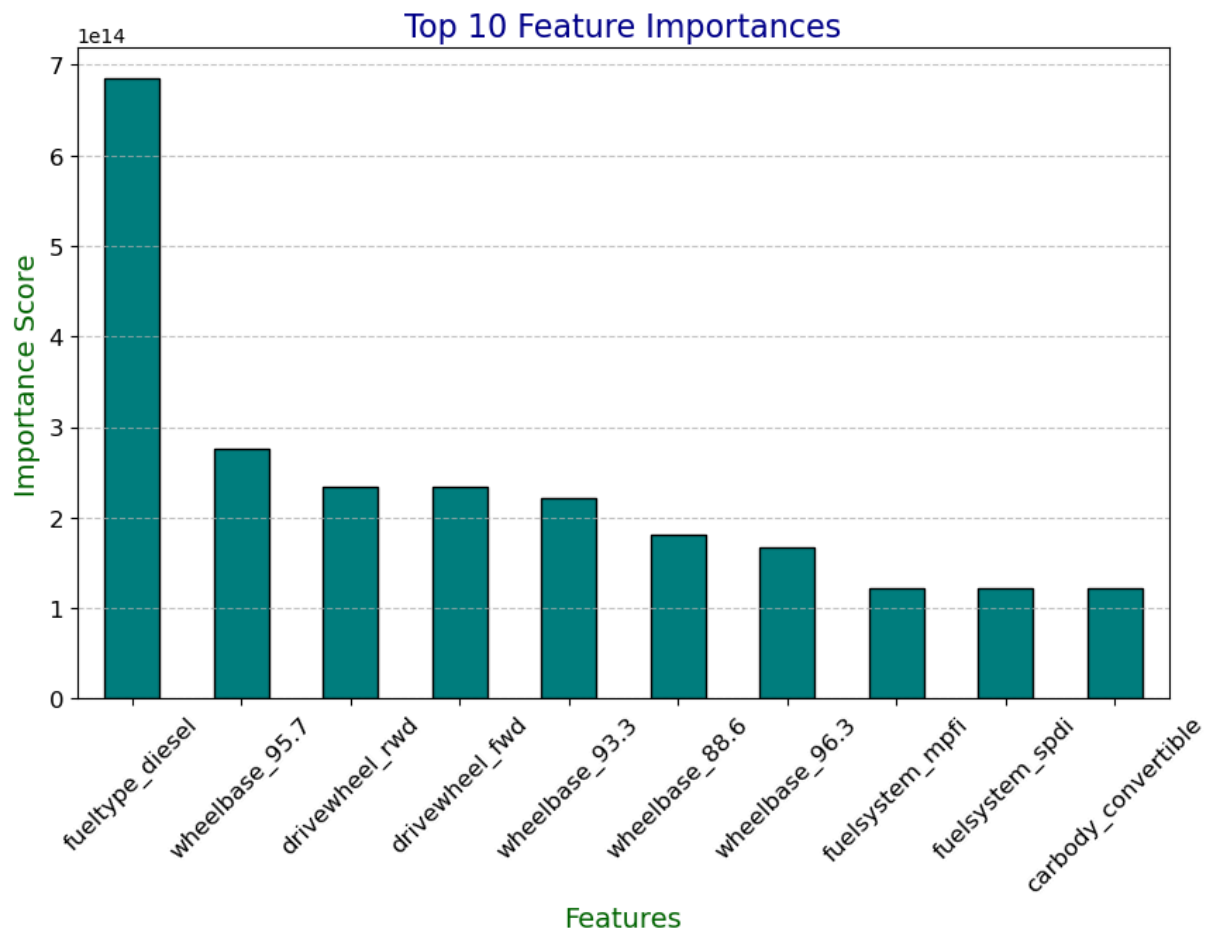
# Display feature importances
print("Feature Importances (from coefficients):\n", feature_importances)
```

Feature Importances (from coefficients):

fueltype_diesel	6.842782e+14
wheelbase_95.7	2.751550e+14
drivewheel_rwd	2.345453e+14
drivewheel_fwd	2.345453e+14
wheelbase_93.3	2.211978e+14
...	
wheelbase_102.0	-4.578681e+14
doornumber_two	-1.114616e+15
doornumber_four	-1.114616e+15
aspiration_turbo	-1.118583e+15
aspiration_std	-1.118583e+15

Length: 554, dtype: float64

```
In [116... plt.figure(figsize=(10, 6))
feature_importances.head(10).plot(kind='bar', color='teal', edgecolor='black')
plt.title('Top 10 Feature Importances', fontsize=16, color='darkblue')
plt.xlabel('Features', fontsize=14, color='darkgreen')
plt.ylabel('Importance Score', fontsize=14, color='darkgreen')
plt.xticks(fontsize=12, rotation=45, color='black')
plt.yticks(fontsize=12, color='black')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```



In [ ]: